



GreenTech Maturity Assessment

Analysis Report

Assessment Details

Date of Assessment

27 December 2024

Project Name

Platform Engineering

Company Name

Tata Consultancy Services (TCS)

Version

1.0

Cloud Provider

AWS

Current Maturity Level

Processes are standardized, documented, well understood and reviewed.

GreenTech Maturity Levels

Maturity Levels

Cover page



Score

Description

Level 1

1

Processes are unpredictable, poorly controlled, and reactive at best.

Level 2

2

Per-project processes. Often still reactive.

Level 3

3

Processes are standardized, documented, well understood and reviewed.

Level 4

4

Processes are measured and controlled.

Level 5

5

Continuous improvement occurs based on quantitative feedback.

Areas Implemented

Phases

Count

Development

Second page

Network

1 of 2

Deployment

7 of 19

Design

0 of 31

Storage

0 of 5

Quality

0 of 8

Operations

0 of 12

Recommendations

Development:

Optimize source code for energy and carbon emissions using static code analysis

Issue Type

Count

Avoid usage of static collections.

87

Avoid multiple if-else statement

61

Do not call a function when declaring a for-type loop

3/29

Use `++i` instead of `i++`

43

String Builder

29

Avoid getting the size of the collection in the loop

23

Avoid creating and starting threads directly

17

Avoid using `Pattern.compile()` in a non-static context

11

Free Resources

2

Use `System.arraycopy` to copy arrays

1

Memory and energy utilization of docker and multi-stage docker files

Metrics

Single-Stage Docker File

Multi-Stage Docker File

Energy Usage

High

Low

Build Time

4/29



5 minutes

6 minutes

Layer Count

12 layers

8 layers

Image Size

2.06 GB

2.01 GB

Cache Efficiency

Low

High

CI/CD Impact

High resource usage

Optimized resource usage

Adopt Multi-Stage Docker Files:

Reduce memory and energy consumption during build and runtime.

Smaller image sizes lead to faster deployments and lower storage needs.

Optimize Base Images:

Use official, slim, or alpine versions of base images to minimize size.

Periodically review and update base images to leverage newer, more efficient versions.

Minimize Dependency Installation:

Only install necessary dependencies to reduce build and runtime resource usage.

Use tools like pip-compile to manage dependencies efficiently.

Leverage Caching:

Utilize Docker's build cache to skip unchanged layers during rebuilds.

Implement caching for dependencies (e.g., using a proxy server for pip dependencies).

Monitor and Analyze Resource Usage:

Employ tools (e.g., Docker Stats, Prometheus, Grafana) to monitor container resource utilization.

Analyze findings to identify optimization opportunities.

Energy metrics for application using Intel RAPL, Kepler, Schaphandre

CPU Model

Intel(R) Xeon(R) Platinum 8259CL CPU @ 2.50GHz

CPU Count

2

GPU Model

N/A

GPU Count

N/A

RAM

3.74 GB

OS

Linux

Country

United States

Region

Virginia



Duration

Energy Consumed

Emissions

Emissions Rate

44.28 Seconds

0.001309 Kwh

0.000483 kg

0.00001091 kg/s

Monitor and Optimize: Regularly collect energy metrics to identify optimization opportunities in your application.

Set Sustainability Goals: Establish targets for your Sustainability Score (SS) or individual energy metrics to drive improvement.

Energy-Aware Scheduling: Utilize Kepler to schedule workflow tasks during periods of low energy demand or when renewable energy sources are available.

Hardware/Software Co-Optimization: Collaborate with hardware teams to optimize system configurations for improved energy efficiency.

Continuously Update and Refine: As new energy-efficient technologies and methodologies emerge, incorporate them into your application and workflows.

Energy consumption of ML job, Model training

Optimize Model Architecture:

Use efficient neural network architectures (e.g., MobileNet, SqueezeNet, EfficientNet) that require less computational resources.

Apply model pruning, knowledge distillation, or quantization to reduce model size and computational requirements.

Select Energy-Efficient Hardware:

Utilize GPU accelerators with high performance-per-watt ratios (e.g., NVIDIA Ampere or AMD CDNA).

Consider TPU (Tensor Processing Unit)-based solutions for large-scale ML workloads.

Explore FPGA (Field-Programmable Gate Array)-based accelerators for customized, energy-efficient computations.

Efficient Training Methodologies:

Employ transfer learning to fine-tune pre-trained models, reducing training time and energy.

Use early stopping techniques to halt training when satisfactory performance is achieved.

Apply distributed training with optimized parallelization to minimize overall training time.

Data Efficiency:

Use smaller, representative datasets for training, reducing computational requirements.

Apply data augmentation to artificially increase dataset size without adding new samples.

Hyperparameter Tuning and Automation:

Utilize hyperparameter tuning tools (e.g., Hyperopt, Optuna) to quickly identify optimal configurations.

Implement automated ML (AutoML) pipelines to streamline the training process and minimize unnecessary computations.

Network:

Energy efficiency for synchronous (REST vs gRPC) and Asynchronous (Kafka vs RabbitMQ) messaging between microservices

Synchronous Communication Protocols

Metrics

gRPC

Rest

Energy Usage

18.1 W

8.2 W

Peak Memory

4.45 MB

8 MB



Time Taken

3 Minutes

4 Min 50 Sec

Requests

500K

500K

Peak CPU

10%

4%

Asynchronous Communication Protocols

Metrics

Kafka

RabbitMQ

Energy Usage

19.2 W

10.68 W

Peak Memory

1024 MB

512 MB

Message Rate

1M msg/s

50K msg/s

Latency

10ms



CPU Usage

55%

35%

For Synchronous Messaging (REST/gRPC)

Use gRPC for New Developments:

Leverage its inherent efficiency advantages, especially for microservices with high inter-service communication.

Migrate REST to HTTP/2:

If moving to gRPC isn't feasible, ensure REST services use HTTP/2 for some efficiency gains.

Optimize Payloads:

Use efficient serialization formats (e.g., protobuf for gRPC, consider alternatives for REST).

Implement compression (if not already done).

Service Discovery and Load Balancing:

Ensure efficient routing to reduce unnecessary network hops.

For Asynchronous Messaging (Kafka/RabbitMQ)

Choose Kafka for High-Volume Scenarios:

Prefer Kafka when dealing with high throughput and low-latency requirements.

Optimize RabbitMQ Configurations:

For existing RabbitMQ setups, review and optimize configurations for batching, compression, and efficient queue management.

Leverage Distributed Capabilities:

For both Kafka and RabbitMQ, ensure distributed setups are optimized to minimize energy consumption.

Deployment:

Tagging of resources to track usage

Mandatory Tags:

Owner (e.g., team, department, individual)

Environment (e.g., dev, staging, prod)

Project/Service (e.g., project name, service identifier)

Cost Center (e.g., budget code, department ID)

Optional Tags (as needed):

Application

Component

Lifecycle (e.g., temporary, permanent)

Compliance (e.g., PCI, HIPAA)

Tagging Best Practices:

Use Meaningful Names: Clearly indicate the tag's purpose.

Keep it Concise: Short tag names (< 20 characters) for easier management.

Avoid Duplication: Use a single tag for a specific attribute (e.g., don't use both Env and Environment).

Use Consistent Formatting: Establish a standard for tag values (e.g., all lowercase, separated by hyphens).

Tagging Structure (Hierarchy):

Flat Structure: Simple, straightforward (e.g., Owner: JohnDoe, Environment: Prod)

Hierarchical Structure: Organized with categories (e.g., Project:MyApp/Environment:Dev, CostCenter:IT/Department:DevOps)

Instances without Tags:

Instance Name

Instance ID

Instance Type



marvel-windows-jumboxtest

i-07a7a6794367ded49

t3.medium

marvel-esha-windows-jumpbox-01

i-0378700bcf2ef68e7

t3.large

marvel-rapl-amd

i-0c43a18a046c87984

t3a.large

marvel-sustain

i-0571933e06349edd6

t2.medium

marvel-autoscaler-3

i-0c86e5cbdc6ef48e9

t2.medium

marvel-docker-registry

i-0121edc6357368099

t3.micro

marvel-git-perforce

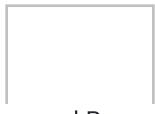
i-0361799692abc3dc4

t2.medium

marvel-esha-windows-jumpbox-03

i-01b006bb7ceefd989

t3.large



marvel-Prometheus

i-00f4b18e0c7615802

t2.micro

marvel-slurm-testing-slurmctl

i-02028574274244586

t3a.small

marvel-slurm-testing-slurm-00

i-0c10cab3a673568a5

t3a.small

marvel-slurm-testing-gw

i-01528621d1601b22d

t3a.small

marvel-slurm-testing-slurmdb

i-04cf6faf1252ba6a7

t3a.small

Calculate energy utilization for provisioned resources and show recommendations to deployment teams

(Legend: R denotes Recommendation Rank)

Under-Provisioned:

Instance ID

Instance Type

R 1

13/29



R 2

R 3

i-0011eee5f9bfa37d5

t3.large

m7i-flex.large

m6i.large

m7i.large

i-01e2bba24c61139a5

t3.large

m7i-flex.large

m6i.large

m7i.large

i-0378700bcf2ef68e7

t3.large

m7i-flex.large

m6i.large

m7i.large

i-061783c19995cf461

t3.large

m7i-flex.large

m6i.large

m7i.large

i-06b0d1885c8bfb73a

t2.medium



c7i-flex.large

c6i.large

c5.large

i-07d971daf49d26ba6

t3.large

m7i-flex.large

m6i.large

m7i.large

i-07fc1239a8432b1ff

t2.nano

t3.small

t3.medium

t3.large

i-0803db063f7902be8

t2.nano

t3.small

t3.medium

t3.large

i-0a30bd8f427568294

t2.large

m7i.large

r7i.large

c7i-flex.xlarge

i-0bcac5f25c482c094



t3.large

m7i-flex.large

m6i.large

m7i.large

i-0e7f58282fba77a57

t3.xlarge

m7i-flex.xlarge

m6i.xlarge

m7i.xlarge

Over-Provisioned:

Instance ID

Instance Type

R 1

R 2

R 3

i-07a7a6794367ded49

t3.xlarge

r7i.large

t3.xlarge

-

i-0f6c16ceb8ea482c2

t3.2xlarge

r6i.xlarge



r7i.xlarge

t3.2xlarge

i-0f97772fef03977df

t3.xlarge

r6i.large

r7i.large

t3.xlarge

i-0fd6069f43c06df1

t3.xlarge

r7i.large

t3.xlarge

-

Identify Unused Resources

Unused EC2 Resources:

Instance Name

Instance ID

Instance Type

Unused Days

marvel-GenAI-partha


i-0a1748bca3199a63f

t3.medium

274

marvel-LLM-test

17/29



i-0b43ba541b65fcd8f

t3.2xlarge

175

marvel-mddp-devsecops

i-0c73c29f8e5d01576

t3.xlarge

162

marvel-EEaaS-Kubernetes

i-0ced2e371b3425fe5

t3.large

142

marvel-TEaaS-openstack-kolla

i-0842262e5bb3e2915

t3.xlarge

141

marvel-awx-ansible

i-0d5da5ba883efcf5b

t3.xlarge

127

marvel-Neureka-registry

i-0dc340b462af4ed74

t3.medium

125

marvel-windows-jumboxtest



i-07a7a6794367ded49

t3.medium

122

marvel-PlatformEng-Keycloak

i-0161cf03311390911

t3.medium

121

marvel-PE-mddp-k8s-1

i-0bfe201057f0db5f9

t3.2xlarge

121

marvel-Prometheus

i-00f4b18e0c7615802

t2.micro

120

marvel-slurm-testing-slurmctl

i-02028574274244586

t3a.small

120

marvel-slurm-testing-slurm-00

i-0c10cab3a673568a5

t3a.small

120

marvel-slurm-testing-gw



i-01528621d1601b22d

t3a.small

120

marvel-slurm-testing-slurmdb

i-04cf6faf1252ba6a7

t3a.small

120

mfdm_qual_dd_pg

i-03b6cc42f6ee8679a

t2.medium

116

tcsnxgnmfmins5

i-054d01dc2082a9968

t3.large

114

tcsnxgnpemonobs01

i-0a30bd8f427568294

t2.large

114

tcsnxgnpemonobs02

i-0644a1225a61cfad0

t3.medium

114

tcsnxgnpemonobs03

20/29



i-0f4a538834c7f8379

t2.large

114

tcsnxgnmfdmins1

i-06f35320739d13110

t3.medium

113

tcsnxgnmfdmins3

i-0dd30488163388001

t3.medium

113

tcsnxgnmfdmins2

i-030fdebbe195c6d05

t3.medium

113

Turn off workloads and node pools outside of business hours

Automate Shutdown/Startup

What It Is: Scripted automatic shutdown/startup of workloads outside business hours.

Sustainability Benefits:

Reduced energy consumption

Lower greenhouse gas emissions

Decreased e-waste from prolonged hardware lifespan

Identify Non-Essential Workloads

What It Is: Analyzing workloads to determine which can be safely turned off during non-business hours.

Sustainability Benefits:

Targeted energy reduction

Minimized unnecessary resource utilization

Enhanced overall efficiency

Autoscaling Configuration (Node Pools)

What It Is: Configuring node pools to autoscale down to 0 nodes during non-business hours.

Sustainability Benefits:

Dynamic energy consumption adjustment

Significant reduction in idle resource energy waste

Enhanced environmental responsiveness

Scheduled Node Pool Management

What It Is: Scheduling node pool shutdowns/startups using platform-specific features or tools.

Sustainability Benefits:

Predictable energy savings

Reduced operational carbon emissions

Improved resource utilization efficiency

Node Pool Sizing Optimization

What It Is: Regular review to optimize node pool sizing.

Sustainability Benefits:

Continuous energy efficiency improvement

Reduced e-waste through minimized node replacements

Lower environmental impact

Show energy efficient resource provisioning options during provisioning

Please refer to this link for a guide on how to implement this.

Implement pre-checks for common issues in CI/CD pipeline to avoid failures in different stages of pipeline

Please refer to this link for a guide on how to implement this.

Details

Development (4/24)

Areas Implemented:

Energy metrics for application using Intel RAPL, Kepler, Schaphandre

Energy consumption of ML job, Model training

Optimize source code for energy and carbon emissions using static code analysis

Compare memory and energy utilization of docker and multi-stage docker files

Areas to be Implemented:

Simplicity and Efficiency Trade-offs while developing algorithms

Efficient Software Algorithms and data structures

Efficient integration and delivery pipelines (Dev, Test, CI env)

On demand development environments

Minimizing unnecessary code execution

Utilizing lazy loading

Optimize energy consumption for Python and Java applications

Leverage LLMs to suggest energy efficient algorithms for Java / Python applications in IDE

Provide configurable deployment of applications with selected features to minimize resource utilization

Calculate energy utilization at function level

Optimize energy utilization of applications using AI/ML

Memory optimization for applications

Rightsizing containers and VM's based on resource utilization

Compare and suggest open-source tools which are energy efficient

Train models with custom rules for optimization of code

Suggest optimized version of built-in functions which are energy efficient in IDE / static analyzer

Code splitting - Convert code to smaller files which can be loaded on-demand

Tree shaking: This technique helps reduce the overall application size by removing unused code from the final build, thus optimizing the web product and minimizing resource consumption

Split large code repositories

Remove dead code and unused data from application

Network (1/2)

Areas Implemented:

Compare energy efficiency for synchronous (REST vs gRPC) and Asynchronous (Kafka vs RabbitMQ) messaging between microservices

Areas to be Implemented:

Optimize data transfer (Implement local data processing, aggregation and compression techniques to avoid transfer of frequent, large amounts of data)

Deployment (7/19)

Areas Implemented:

Tagging of resources to track usage

Calculate energy utilization for provisioned resources and show recommendations to deployment teams

Identify and share report on unused resources

Turn off workloads and node pools outside of business hours

Show energy efficient resource provisioning options during provisioning

Implement pre-checks for common issues in CI/CD pipeline to avoid failures in different stages of pipeline

Workload allocation on energy efficient servers, VMs

Areas to be Implemented:

Deploy application in a region that is closest to users

Deploy in regions that are powered by renewables or energy efficient resources

Energy efficient workload scheduling

CPU frequency optimization for non-critical workloads

Alerts during resource provisioning (to avoid over provisioning) about energy consumption and energy efficient options

Calculate and minimize energy utilization per transaction

Implement automated decision between horizontal vs vertical scaling of applications

Set up minimal artifacts - Reduce disk, memory and processing demands by considering the software for a given purpose

Schedule resource intensive tasks to execute when renewable energy source is available

Consolidate application workloads to maximize server utilization and move idle servers to low power mode

Measure (Idle/Normal/Peak) energy consumption of resources and identify scope for optimizations (Get power consumption per CPU core information from processor type)

Identify unused resources in K8s using KOR tool

Design (0/31)

Areas Implemented:

No Areas Implemented

Areas to be Implemented:

Sustainability as NFR

Minimal Architecture (i.e. Developing minimal software for value)

Replace long-running service with a simple Function as a Service

Asynchronous processing and event-driven architectures

API Design: protocol (TCP, UDP, custom wire protocols)

API Design: data format used for request/response of your APIs: Protobuf/gRPC is better compared to JSON, XML

API Design: Avro schemas instead of JSON for data serialization

Reusable APIs vs P2P Integrations

Choice of programming language: Go, Rust over Python, Java; Java 17 is better with respect to memory optimization

Public cloud, choice of provider, and region

Platform as a Service (PaaS) and Serverless (e.g., Function as a Service (FaaS))

Containers and Kubernetes - Adopting energy efficient architectures such as containers or serverless

Scheduling and batch vs. real-time - Demand Shaping

Carbon awareness into the Kubernetes Scheduler

Cost and Sustainability: Reducing operational costs and sustainability are aligned

Enforce Quotas and Rate Limiting

Reduce the Network Footprint

Identify energy utilization of applications across different design choices

Optimize data generation of applications

Create reusable modules to minimize development and maintenance efforts

Enable parallel processing using distributed computing and minimize energy consumption

Measure KPIs on the features used by customers and remove unused features

Energy saving mode for applications, system processes during periods of low load

Energy budgets for applications

Create Energy profile for application (Core Vs Non-core tasks - Logging, Monitoring, Audit, Redundancy etc...)

Improve resilience of equipment to minimize redundancy (Ensure legacy hardware can support at least core features of software)

Shared libraries for applications

Server-side rendering (SSR): With SSR, a web page's HTML is generated on the server and sent to the client, resulting in quicker initial load times and improved search engine optimization

React frameworks: Next.js as a framework has many built-in features for energy optimization

Progressive Web Applications (Hybrid of native and web apps) are more energy efficient

Optimize availability and scalability needs for application

Storage (0/5)

Areas Implemented:

No Areas Implemented

Areas to be Implemented:

Dynamic provisioning of volumes for EC2 instance / Pod to automatically scale the storage size

Identify and remove unused data and resources in cloud

Implementing data retention, archival and deletion policies

Utilize caching to store frequently used data

Optimize log and metrics collection

Quality (0/8)

Areas Implemented:

No Areas Implemented

Areas to be Implemented:

Code and design review for sustainability

Document decisions using Key Design Decisions (KDDs) or Architectural Decision Records (ADRs)

Profiling tools and static analysis - Code profiling and optimization tools

Upgrade Runtimes and Modules

Change aware testing

Test Case Optimization

Risk Based Testing

On-demand and scalable Test Environments

Operations (0/12)

Areas Implemented:

No Areas Implemented

Areas to be Implemented:

Automated sustainability maturity assessment for applications

Observability and CarbonOps

Automation

Switching instances to chipsets that offer the same processing power at lower levels of energy consumption

Sustainable DevOps practices integrate sustainability into the software delivery process

Optimize resource allocation using ML

Maximize utilization of resources

Spot Instances

Setup sustainability goals for infrastructure, development and applications and monitor them using dashboards and reports

Reduce the cooling requirements of servers by ensuring that applications (servers) run in optimal temperature range

Identify zombie workloads that consume resources (e.g., Application not running within container / VM)

Dashboard for resource utilization at cluster, node, and pod level and grouped by applications in Kubernetes