GreenTech Maturity Assessment

Analysis Report

Assessment Details

Date of Assessment	**27 December 2024**

Cover page

Project Name	**Platform Engineering**
Company Name	**Tata Consultancy Services (TCS)**
Version	**1.0**
Cloud Provider	**AWS**

Current Maturity Level

Processes are standardized, documented, well understood and reviewed.

Maturity Level
03

GreenTech Maturity Levels

Leve l 1	**1**	**Processes are unpredictable, poorly controlled, and reactive at best.**

Second page

Level 2	2	Per-project processes. Often still reactive.
Level 3	3	Processes are standardized, documented, well understood and reviewed.
Level 4	4	Processes are measured and controlled.
Level 5	5	Continuous improvement occurs based on quantitative feedback.

Areas Implemented

Development	4 of 24
Network	1 of 2

Deployment	**7 of 19**
Design	**0 of 31**
Storage	**0 of 5**
Quality	**0 of 8**
Operations	**0 of 12**

Recommendations

Development:

1. **Optimize source code for energy and carbon emissions using static code analysis**

Avoid usage of static collections.	**87**
Avoid multiple if-else statement	**61**
Do not call a function when declaring a for-type loop	**54**

Use ++i instead of i++	43
String Builder	29
Avoid getting the size of the collection in the loop	23
Avoid creating and starting threads directly	17
Avoid using Pattern.compile() in a non-static context	11
Free Resources	2
Use System. arraycopy to copy arrays	1

2. **Memory and energy utilization of docker and multi-stage docker files**

Energy Usage	High	Low
Build Time	5 minutes	6 minutes

Layer Count	12 layers	8 layers
Image Size	2.06 GB	2.01 GB
Cache Efficiency	Low	High
CI/CD Impact	High resource usage	Optimized resource usage

Adopt Multi-Stage Docker Files:

- Reduce memory and energy consumption during build and runtime.

- Smaller image sizes lead to faster deployments and lower storage needs.

Optimize Base Images:

- Use official, slim, or alpine versions of base images to minimize size.

- Periodically review and update base images to leverage newer, more efficient versions.

Minimize Dependency Installation:

- Only install necessary dependencies to reduce build and runtime resource usage.

- Use tools like pip-compile to manage dependencies efficiently.

Leverage Caching:

- Utilize Docker's build cache to skip unchanged layers during rebuilds.

- Implement caching for dependencies (e.g., using a proxy server for pip dependencies).

Monitor and Analyze Resource Usage:

- Employ tools (e.g., Docker Stats, Prometheus, Grafana) to monitor container resource utilization.

- Analyze findings to identify optimization opportunities.

3. **Energy metrics for application using Intel RAPL, Kepler, Schaphandre**

CPU Model	Intel(R) Xeon(R) Platinum 8259CL CPU @ 2.50GHz
CPU Count	2
GPU Model	N/A
GPU Count	N/A
RAM	3.74 GB

OS	Linux
Country	United States
Region	Virginia

44.28 Seconds	0.001309 Kwh	0.000483 kg	0.00001091 kg/s

Monitor and Optimize: Regularly collect energy metrics to identify optimization opportunities in your application.

Set Sustainability Goals: Establish targets for your Sustainability Score (SS) or individual energy metrics to drive improvement.

Energy-Aware Scheduling: Utilize Kepler to schedule workflow tasks during periods of low energy demand or when renewable energy sources are available.

Hardware/Software Co-Optimization: Collaborate with hardware teams to optimize system configurations for improved energy efficiency.

Continuously Update and Refine: As new energy-efficient technologies and methodologies emerge, incorporate them into your application and workflows.

4. **Energy consumption of ML job, Model training**

Optimize Model Architecture:

- Use efficient neural network architectures (e.g., MobileNet, SqueezeNet, EfficientNet) that require less computational resources.

- Apply model pruning, knowledge distillation, or quantization to reduce model size and computational requirements.

Select Energy-Efficient Hardware:

- Utilize GPU accelerators with high performance-per-watt ratios (e.g., NVIDIA Ampere or AMD CDNA).
- Consider TPU (Tensor Processing Unit)-based solutions for large-scale ML workloads.
- Explore FPGA (Field-Programmable Gate Array)-based accelerators for customized, energy-efficient computations.

Efficient Training Methodologies:

- Employ transfer learning to fine-tune pre-trained models, reducing training time and energy.
- Use early stopping techniques to halt training when satisfactory performance is achieved.
- Apply distributed training with optimized parallelization to minimize overall training time.

Data Efficiency:

- Use smaller, representative datasets for training, reducing computational requirements.
- Apply data augmentation to artificially increase dataset size without adding new samples.

Hyperparameter Tuning and Automation:

- Utilize hyperparameter tuning tools (e.g., Hyperopt, Optuna) to quickly identify optimal configurations.

- Implement automated ML (AutoML) pipelines to streamline the training process and minimize unnecessary computations.

Network:

1. **Energy efficiency for synchronous (REST vs gRPC) and Asynchronous (Kafka vs RabbitMQ) messaging between microservices**

Synchronous Communication Protocols

Energy Usage	18.1 W	8.2 W
Peak Memory	4.45 MB	8 MB
Time Taken	3 Minutes	4 Min 50 Sec
Requests	500K	500K
Peak CPU	10%	4%

Asynchronous Communication Protocols

Energy Usage	19.2 W	10.68 W
Peak Memory	1024 MB	512 MB
Message Rate	1M msg/s	50K msg/s
Latency	10ms	5ms
CPU Usage	55%	35%

Energy Usage (Watts) chart showing gRPC ~18, REST ~8, Kafka ~19, RabbitMQ ~10.5

For Synchronous Messaging (REST/gRPC)

Use gRPC for New Developments:

- Leverage its inherent efficiency advantages, especially for microservices with high inter-service communication.

Migrate REST to HTTP/2:

- If moving to gRPC isn't feasible, ensure REST services use HTTP/2 for some efficiency gains.

Optimize Payloads:

- Use efficient serialization formats (e.g., protobuf for gRPC, consider alternatives for REST).

- Implement compression (if not already done).

Service Discovery and Load Balancing:

- Ensure efficient routing to reduce unnecessary network hops.

For Asynchronous Messaging (Kafka/RabbitMQ)

Choose Kafka for High-Volume Scenarios:

- Prefer Kafka when dealing with high throughput and low-latency requirements.

Optimize RabbitMQ Configurations:

- For existing RabbitMQ setups, review and optimize configurations for batching, compression, and efficient queue management.

Leverage Distributed Capabilities:

- For both Kafka and RabbitMQ, ensure distributed setups are optimized to minimize energy consumption.

Deployment:

1. **Tagging of resources to track usage**

Mandatory Tags:

- Owner (e.g., team, department, individual)

- Environment (e.g., dev, staging, prod)

- Project/Service (e.g., project name, service identifier)

- Cost Center (e.g., budget code, department ID)

Optional Tags (as needed):

- Application

- Component

- Lifecycle (e.g., temporary, permanent)

- Compliance (e.g., PCI, HIPAA)

Tagging Best Practices:

- Use Meaningful Names: Clearly indicate the tag's purpose.

- Keep it Concise: Short tag names (< 20 characters) for easier management.

- Avoid Duplication: Use a single tag for a specific attribute (e.g., don't use both Env and Environment).

- Use Consistent Formatting: Establish a standard for tag values (e.g., all lowercase, separated by hyphens).

Tagging Structure (Hierarchy):

- Flat Structure: Simple, straightforward (e.g., Owner: JohnDoe, Environment: Prod)

- Hierarchical Structure: Organized with categories (e.g., Project:MyApp/Environment:Dev, CostCenter:IT/Department:DevOps)

Instances without Tags:

marvel-windows-jumboxtest	**i-07a7a6794367ded49**	**t3.medium**
marvel-esha-windows-jumpbox-01	**i-0378700bcf2ef68e7**	**t3.large**
marvel-rapl-amd	**i-0c43a18a046c87984**	**t3a.large**
marvel-sustain	**i-0571933e06349edd6**	**t2.medium**
marvel-autoscaler-3	**i-0c86e5cbdc6ef48e9**	**t2.medium**
marvel-docker-registry	**i-0121edc6357368099**	**t3.micro**
marvel-git-perforce	**i-0361799692abc3dc4**	**t2.medium**
marvel-esha-windows-jumpbox-03	**i-01b006bb7ceefd989**	**t3.large**

marvel-Prometheus	**i-00f4b18e0c7615802**	**t2.micro**
marvel-slurm-testing-slurmctl	**i-02028574274244586**	**t3a.small**
marvel-slurm-testing-slurm-00	**i-0c10cab3a673568a5**	**t3a.small**
marvel-slurm-testing-gw	**i-01528621d1601b22d**	**t3a.small**
marvel-slurm-testing-slurmdb	**i-04cf6faf1252ba6a7**	**t3a.small**

2. **Calculate energy utilization for provisioned resources and show recommendations to deployment teams**

(Legend: R denotes Recommendation Rank)

Under-Provisioned:

i-0011eee5f9bfa37d5	**t3.large**	**m7i-flex.large**	**m6i.large**	**m7i.large**

i-01e2bba24c61139a5	**t3.large**	**m7i-flex.large**	**m6i.large**	**m7i.large**
i-0378700bcf2ef68e7	**t3.large**	**m7i-flex.large**	**m6i.large**	**m7i.large**
i-061783c19995cf461	**t3.large**	**m7i-flex.large**	**m6i.large**	**m7i.large**
i-06b0d1885c8bfb73a	**t2.medium**	**c7i-flex.large**	**c6i.large**	**c5.large**
i-07d971daf49d26ba6	**t3.large**	**m7i-flex.large**	**m6i.large**	**m7i.large**
i-07fc1239a8432b1ff	**t2.nano**	**t3.small**	**t3.medium**	**t3.large**
i-0803db063f7902be8	**t2.nano**	**t3.small**	**t3.medium**	**t3.large**
i-0a30bd8f427568294	**t2.large**	**m7i.large**	**r7i.large**	**c7i-flex.xlarge**
i-0bcac5f25c482c094	**t3.large**	**m7i-flex.large**	**m6i.large**	**m7i.large**
i-0e7f58282fba77a57	**t3.xlarge**	**m7i-flex.xlarge**	**m6i.xlarge**	**m7i.xlarge**

Over-Provisioned:

i-07a7a6794367ded49	t3.xlarge	r7i.large	t3.xlarge	-
i-0f6c16ceb8ea482c2	t3.2xlarge	r6i.xlarge	r7i.xlarge	t3.2xlarge
i-0f97772fef03977df	t3.xlarge	r6i.large	r7i.large	t3.xlarge
i-0fda6069f43c06df1	t3.xlarge	r7i.large	t3.xlarge	-

3. **Identify Unused Resources**

Unused EC2 Resources:

marvel-GenAI-partha	i-0a1748bca3199a63f	t3.medium	274
marvel-LLM-test	i-0b43ba541b65fcd8f	t3.2xlarge	175
marvel-mddp-devsecops	i-0c73c29f8e5d01576	t3.xlarge	162
marvel-EEaaS-Kubernetes	i-0ced2e371b3425fe5	t3.large	142

marvel-TEaas-openstack-kolla	**i-0842262e5bb3e2915**	**t3.xlarge**	**141**
marvel-awx-ansible	**i-0d5da5ba883efcf5b**	**t3.xlarge**	**127**
marvel-Neureda-registry	**i-0dc340b462af4ed74**	**t3.medium**	**125**
marvel-windows-jumboxtest	**i-07a7a6794367ded49**	**t3.medium**	**122**
marvel-PlatformEng-Keycloak	i-0161cf03311390911	**t3.medium**	**121**
marvel-PE-mddp-k8s-1	**i-0bfe201057f0db5f9**	**t3.2xlarge**	**121**
marvel-Prometheus	**i-00f4b18e0c7615802**	**t2.micro**	**120**
marvel-slurm-testing-slurmctl	**i-02028574274244586**	**t3a.small**	**120**
marvel-slurm-testing-slurm-00	**i-0c10cab3a673568a5**	**t3a.small**	**120**
marvel-slurm-testing-gw	**i-01528621d1601b22d**	**t3a.small**	**120**
marvel-slurm-testing-slurmdb	**i-04cf6faf1252ba6a7**	**t3a.small**	**120**
mfdm_qual_dd_pg	**i-03b6cc42f6ee8679a**	**t2.medium**	**116**

tcsnxgnmfdmins5	**i-054d01dc2082a9968**	**t3.large**	**114**
tcsnxgnpemonobs01	**i-0a30bd8f427568294**	**t2.large**	**114**
tcsnxgnpemonobs02	**i-0644a1225a61cfad0**	**t3.medium**	**114**
tcsnxgnpemonobs03	**i-0f4a538834c7f8379**	**t2.large**	**114**
tcsnxgnmfdmins1	**i-06f35320739d13110**	**t3.medium**	**113**
tcsnxgnmfdmins3	**i-0dd30488163388001**	**t3.medium**	**113**
tcsnxgnmfdmins2	**i-030fdebbe195c6d05**	**t3.medium**	**113**

4. **Turn off workloads and node pools outside of business hours**

Automate Shutdown/Startup

- What It Is: Scripted automatic shutdown/startup of workloads outside business hours.

Sustainability Benefits:

- Reduced energy consumption

- Lower greenhouse gas emissions

- Decreased e-waste from prolonged hardware lifespan

Identify Non-Essential Workloads

- What It Is: Analyzing workloads to determine which can be safely turned off during non-business hours.

Sustainability Benefits:

- Targeted energy reduction

- Minimized unnecessary resource utilization

- Enhanced overall efficiency

Autoscaling Configuration (Node Pools)

- What It Is: Configuring node pools to autoscale down to 0 nodes during non-business hours.

Sustainability Benefits:

- Dynamic energy consumption adjustment

- Significant reduction in idle resource energy waste

- Enhanced environmental responsiveness

Scheduled Node Pool Management

- What It Is: Scheduling node pool shutdowns/startups using platform-specific features or tools.

Sustainability Benefits:

- Predictable energy savings
- Reduced operational carbon emissions
- Improved resource utilization efficiency

Node Pool Sizing Optimization

- What It Is: Regular review to optimize node pool sizing.

Sustainability Benefits :

- Continuous energy efficiency improvement
- Reduced e-waste through minimized node replacements
- Lower environmental impact

5. **Show energy efficient resource provisioning options during provisioning**

Please refer to this [link](#) for a guide on how to implement this.

6. **Implement pre-checks for common issues in CI/CD pipeline to avoid failures in different stages of pipeline**

Please refer to this [link](#) for a guide on how to implement this.

Details

1. **Development (4/24)**

Areas Implemented:

- Energy metrics for application using Intel RAPL, Kepler, Schaphandre

- Energy consumption of ML job, Model training

- Optimize source code for energy and carbon emissions using static code analysis

- Compare memory and energy utilization of docker and multi-stage docker files

Areas to be Implemented:

- Simplicity and Efficiency Trade-offs while developing algorithms

- Efficient Software Algorithms and data structures

- Efficient integration and delivery pipelines (Dev, Test, CI env)

- On demand development environments

- Minimizing unnecessary code execution

- Utilizing lazy loading

- Optimize energy consumption for Python and Java applications

- Leverage LLMs to suggest energy efficient algorithms for Java / Python applications in IDE

- Provide configurable deployment of applications with selected features to minimize resource utilization

- Calculate energy utilization at function level

- Optimize energy utilization of applications using AI/ML

- Memory optimization for applications

- Rightsizing containers and VM's based on resource utilization

- Compare and suggest open-source tools which are energy efficient

- Train models with custom rules for optimization of code

- Suggest optimized version of built-in functions which are energy efficient in IDE / static analyzer

- Code splitting - Convert code to smaller files which can be loaded on-demand
- Tree shaking: This technique helps reduce the overall application size by removing unused code from the final build, thus optimizing the web product and minimizing resource consumption
- Split large code repositories
- Remove dead code and unused data from application

2. **Network (1/2)**

Areas Implemented:

- Compare energy efficiency for synchronous (REST vs gRPC) and Asynchronous (Kafka vs RabbitMQ) messaging between microservices

Areas to be Implemented:

- Optimize data transfer (Implement local data processing, aggregation and compression techniques to avoid transfer of frequent, large amounts data)

3. **Deployment (7/19)**

Areas Implemented:

- Tagging of resources to track usage
- Calculate energy utilization for provisioned resources and show recommendations to deployment teams
- Identify and share report on unused resources
- Turn off workloads and node pools outside of business hours
- Show energy efficient resource provisioning options during provisioning
- Implement pre-checks for common issues in CI/CD pipeline to avoid failures in different stages of pipeline
- Workload allocation on energy efficient servers, VMs

Areas to be Implemented:

- Deploy application in a region that is closest to users
- Deploy in regions that are powered by renewables or energy efficient resources
- Energy efficient workload scheduling
- CPU frequency optimization for non-critical workloads

- Alerts during resource provisioning (to avoid over provisioning) about energy consumption and energy efficient options

- Calculate and minimize energy utilization per transaction

- Implement automated decision between horizontal vs vertical scaling of applications

- Set up minimal artifacts - Reduce disk, memory and processing demands by considering the software for a given purpose

- Schedule resource intensive tasks to execute when renewable energy source is available

- Consolidate application workloads to maximize server utilization and move idle servers to low power mode

- Measure (Idle/Normal/Peak) energy consumption of resources and identify scope for optimizations (Get power consumption per CPU core information from processor type)

- Identify unused resources in K8s using KOR tool

4. **Design (0/31)**

Areas Implemented:

- No Areas Implemented

Areas to be Implemented:

- Sustainability as NFR

- Minimal Architecture (i.e. Developing minimal software for value)

- Replace long-running service with a simple Function as a Service

- Asynchronous processing and event-driven architectures

- API Design: protocol (TCP, UDP, custom wire protocols)

- API Design: data format used for request/response of your APIs: Protobuf/gRPC is better compared to JSON, XML

- API Design: Avro schemas instead of JSON for data serialization

- Reusable APIs vs P2P Integrations

- Choice of programming language: Go, Rust over Python, Java; Java 17 is better with respect to memory optimization

- Public cloud, choice of provider, and region

- Platform as a Service (PaaS) and Serverless (e.g., Function as a Service (FaaS))

- Containers and Kubernetes - Adopting energy efficient architectures such as containers or serverless

- Scheduling and batch vs. real-time - Demand Shaping

- Carbon awareness into the Kubernetes Scheduler

- Cost and Sustainability: Reducing operational costs and sustainability are aligned

- Enforce Quotas and Rate Limiting

- Reduce the Network Footprint

- Identify energy utilization of applications across different design choices

- Optimize data generation of applications

- Create reusable modules to minimize development and maintenance efforts

- Enable parallel processing using distributed computing and minimize energy consumption

- Measure KPIs on the features used by customers and remove unused features

- Energy saving mode for applications, system processes during periods of low load

- Energy budgets for applications

- Create Energy profile for application (Core Vs Non-core tasks - Logging, Monitoring, Audit, Redundancy etc...)

- Improve resilience of equipment to minimize redundancy (Ensure legacy hardware can support at least core features of software)

- Shared libraries for applications

- Server-side rendering (SSR): With SSR, a web page's HTML is generated on the server and sent to the client, resulting in quicker initial load tir and improved search engine optimization

- React frameworks: Next.js as a framework has many built-in features for energy optimization

- Progressive Web Applications (Hybrid of native and web apps) are more energy efficient

- Optimize availability and scalability needs for application

5. **Storage (0/5)**

Areas Implemented:

- No Areas Implemented

Areas to be Implemented:

- Dynamic provisioning of volumes for EC2 instance / Pod to automatically scale the storage size
- Identify and remove unused data and resources in cloud
- Implementing data retention, archival and deletion policies
- Utilize caching to store frequently used data

- Optimize log and metrics collection

6. **Quality (0/8)**

Areas Implemented:

- No Areas Implemented

Areas to be Implemented:

- Code and design review for sustainability
- Document decisions using Key Design Decisions (KDDs) or Architectural Decision Records (ADRs)
- Profiling tools and static analysis - Code profiling and optimization tools
- Upgrade Runtimes and Modules
- Change aware testing
- Test Case Optimization
- Risk Based Testing

- On-demand and scalable Test Environments

7. **Operations (0/12)**

Areas Implemented:

- No Areas Implemented

Areas to be Implemented:

- Automated sustainability maturity assessment for applications
- Observability and CarbonOps
- Automation
- Switching instances to chipsets that offer the same processing power at lower levels of energy consumption
- Sustainable DevOps practices integrate sustainability into the software delivery process
- Optimize resource allocation using ML
- Maximize utilization of resources

- Spot Instances

- Setup sustainability goals for infrastructure, development and applications and monitor them using dashboards and reports

- Reduce the cooling requirements of servers by ensuring that applications (servers) run in optimal temperature range

- Identify zombie workloads that consume resources (e.g., Application not running within container / VM)

- Dashboard for resource utilization at cluster, node, and pod level and grouped by applications in Kubernetes

Last Page