GreenTech Maturity Assessment

Analysis Report

Assessment Details :

Date of Assessment	27 December 2024
Project Name	Platform Engineering
Company Name	Tata Consultancy Services (TCS)
Version	1.0
Cloud Provider	AWS

Current Maturity Level :
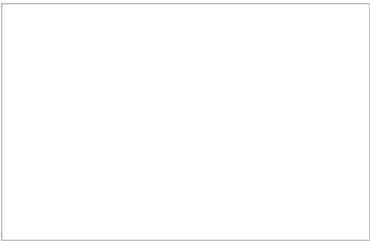
Processes are standardized, documented, well understood and reviewed.

GreenTech Maturity Levels :

Maturity Levels	Score	Description

Cover page

Level 1	1	Processes are unpredictable, poorly controlled, and reactive at best.
Level 2	2	Per-project processes. Often still reactive.
Level 3	3	Processes are standardized, documented, well understood and reviewed.
Level 4	4	Processes are measured and controlled.
Level 5	5	Continuous improvement occurs based on quantitative feedback.

Areas Implemented :

Phases	Count
Development	4 of 24
Network	1 of 2
Deployment	7 of 19
Design	0 of 31
Storage	0 of 5
Quality	0 of 8
Operations	0 of 12

Recommendations :

Development:

1. Optimize source code for energy and carbon emissions using static code analysis

Issue Type	Count
Avoid usage of static collections.	87
Avoid multiple if-else statement	61
Do not call a function when declaring a for-type loop	54
Use ++i instead of i++	43
String Builder	29
Avoid getting the size of the collection in the loop	23

Second page

Avoid creating and starting threads directly	*17*
Avoid using Pattern.compile() in a non-static context	*11*
Free Resources	*2*
Use System.arraycopy to copy arrays	*1*

2. Memory and energy utilization of docker and multi-stage docker files

Metrics	Single-Stage Docker File	Multi-Stage Docker File
Energy Usage	*High*	*Low*
Build Time	*5 minutes*	*6 minutes*
Layer Count	*12 layers*	*8 layers*
Image Size	*2.06 GB*	*2.01 GB*
Cache Efficiency	*Low*	*High*
CI/CD Impact	*High resource usage*	*Optimized resource usage*

Adopt Multi-Stage Docker Files :

- *Reduce memory and energy consumption during build and runtime.*
- *Smaller image sizes lead to faster deployments and lower storage needs.*

Optimize Base Images :

- *Use official, slim, or alpine versions of base images to minimize size.*
- *Periodically review and update base images to leverage newer, more efficient versions.*

Minimize Dependency Installation :

- *Only install necessary dependencies to reduce build and runtime resource usage.*
- *Use tools like pip-compile to manage dependencies efficiently.*

Leverage Caching :

- *Utilize Docker's build cache to skip unchanged layers during rebuilds.*
- *Implement caching for dependencies (e.g., using a proxy server for pip dependencies).*

Monitor and Analyze Resource Usage :

- *Employ tools (e.g., Docker Stats, Prometheus, Grafana) to monitor container resource utilization.*
- *Analyze findings to identify optimization opportunities.*

3. Energy metrics for application using Intel RAPL, Kepler, Schaphandre

CPU Model	Intel(R) Xeon(R) Platinum 8259CL CPU @ 2.50GHz
CPU Count	2
GPU Model	N/A
GPU Count	N/A
RAM	3.74 GB
OS	Linux
Country	United States
Region	Virginia

Duration	Energy Consumed	Emissions	Emissions Rate
44.28 Seconds	0.001309 Kwh	0.000483 kg	0.00001091 kg/s
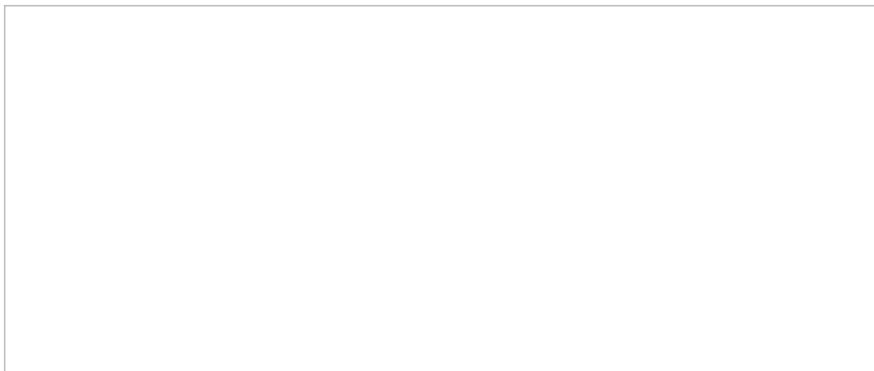
Monitor and Optimize :   Regularly collect energy metrics to identify optimization opportunities in your application.

Set Sustainability Goals :   Establish targets for your Sustainability Score (SS) or individual energy metrics to drive improvement.

Energy-Aware Scheduling :   Utilize Kepler to schedule workflow tasks during periods of low energy demand or when renewable energy sources are available.

Continuously Update and Refine :   Collaborate with hardware teams to optimize system configurations for improved energy efficiency.

Hardware/Software Co-Optimization :   As new energy-efficient technologies and methodologies emerge, incorporate them into your application and workflows.

4. Energy consumption of ML job, Model training

Optimize Model Architecture:

- Use efficient neural network architectures (e.g., MobileNet, SqueezeNet, EfficientNet) that require less

computational resources.

- Apply model pruning, knowledge distillation, or quantization to reduce model size and computational requirements.

Select Energy-Efficient Hardware :

- Utilize GPU accelerators with high performance-per-watt ratios (e.g., NVIDIA Ampere or AMD CDNA).
- Consider TPU (Tensor Processing Unit)-based solutions for large-scale ML workloads.
- Explore FPGA (Field-Programmable Gate Array)-based accelerators for customized, energy-efficient computations.

Efficient Training Methodologies :

- Employ transfer learning to fine-tune pre-trained models, reducing training time and energy.
- Use early stopping techniques to halt training when satisfactory performance is achieved.
- Apply distributed training with optimized parallelization to minimize overall training time.

Data Efficiency:

- Use smaller, representative datasets for training, reducing computational requirements.
- Apply data augmentation to artificially increase dataset size without adding new samples.

Hyperparameter Tuning and Automation :

- Utilize hyperparameter tuning tools (e.g., Hyperopt, Optuna) to quickly identify optimal configurations.
- Implement automated ML (AutoML) pipelines to streamline the training process and minimize unnecessary computations.

Network :

1. Energy efficiency for synchronous (REST vs gRPC) and Asynchronous (Kafka vs RabbitMQ) messaging between microservices
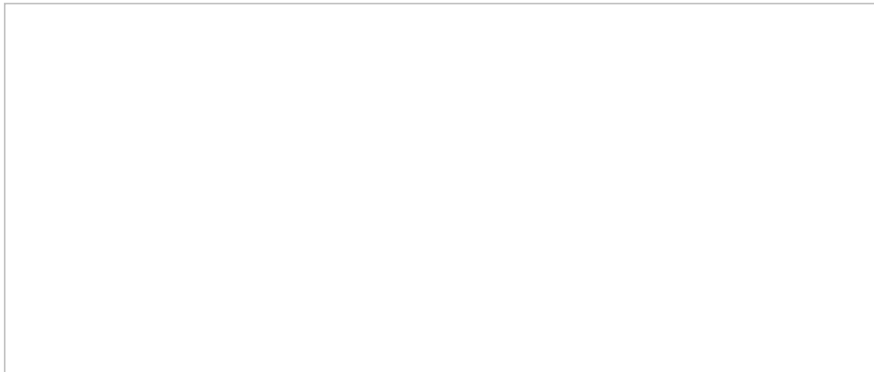
Synchronous Communication Protocols

Metrics	gRPC	Rest
Energy Usage	18.1 W	8.2 W
Peak Memory	4.45 MB	8 MB
Time Taken	3 Minutes	4 Min 50 Sec
Requests	500K	500K
Peak CPU	10%	4%

Asynchronous Communication Protocols

Metrics	Kafka	RabbitMQ

Energy Usage	19.2 W	10.68 W
Peak Memory	1024 MB	512 MB
Message Rate	1M msg/s	50K msg/s
Latency	10ms	5ms
CPU Usage	55%	35%

For Synchronous Messaging (REST/gRPC)

Use gRPC for New Developments :

- Leverage its inherent efficiency advantages, especially for microservices with high inter-service communication.

Migrate REST to HTTP/2 :

- If moving to gRPC isn't feasible, ensure REST services use HTTP/2 for some efficiency gains.

Optimize Payloads :

- Use efficient serialization formats (e.g., protobuf for gRPC, consider alternatives for REST).
- Implement compression (if not already done).

Service Discovery and Load Balancing :

- Ensure efficient routing to reduce unnecessary network hops.

For Asynchronous Messaging (Kafka/RabbitMQ)

Use gRPC for New Developments :

- Leverage its inherent efficiency advantages, especially for microservices with high inter-service communication.

Migrate REST to HTTP/2 :

- If moving to gRPC isn't feasible, ensure REST services use HTTP/2 for some efficiency gains.

Optimize Payloads :

- *Use efficient serialization formats (e.g., protobuf for gRPC, consider alternatives for REST).*
- *Implement compression (if not already done).*

Service Discovery and Load Balancing :

- *Ensure efficient routing to reduce unnecessary network hops.*

Deployment :

1. Tagging of resources to track usage

Mandatory Tags :

- *Owner (e.g., team, department, individual)*
- *Environment (e.g., dev, staging, prod)*
- *Project/Service (e.g., project name, service identifier)*
- *Cost Center (e.g., budget code, department ID)*

Optional Tags (as needed) :

- *Application*
- *Component*
- *Lifecycle (e.g., temporary, permanent)*
- *Compliance (e.g., PCI, HIPAA)*

Tagging Best Practices :

- *Use Meaningful Names: Clearly indicate the tag's purpose.*
- *Keep it Concise: Short tag names (< 20 characters) for easier management.*
- *Avoid Duplication: Use a single tag for a specific attribute (e.g., don't use both Env and Environment).*
- *Use Consistent Formatting: Establish a standard for tag values (e.g., all lowercase, separated by hyphens).*

Tagging Structure (Hierarchy) :

- *Flat Structure: Simple, straightforward (e.g., Owner: JohnDoe, Environment: Prod)*
- *Hierarchical Structure: Organized with categories (e.g., Project:MyApp/Environment:Dev, CostCenter:IT/Department:DevOps)*

Instances without Tags:

Instance Name	Instance Name	Instance Type
marvel-windows-jumboxtest	*i-07a7a6794367ded49*	*t3.medium*
marvel-esha-windows-jumpbox-01	*i-0378700bcf2ef68e7*	*t3.large*
marvel-rapl-amd	*i-0c43a18a046c87984*	*t3a.large*
marvel-sustain	*i-0571933e06349edd6*	*t2.medium*
marvel-autoscaler-3	*i-0c86e5cbdc6ef48e9*	*t2.medium*

marvel-esha-windows-jumpbox-03	*i-01b006bb7ceefd989*	*t3.large*
marvel-docker-registry	*i-0121edc6357368099*	*t3.micro*
marvel-git-perforce	*i-0361799692abc3dc4*	*t2.medium*
marvel-esha-windows-jumpbox-03	*i-01b006bb7ceefd989*	*t3.large*
marvel-Prometheus	*i-00f4b18e0c7615802*	*t2.micro*
marvel-slurm-testing-slurmctl	*i-02028574274244586*	*t3a.small*
marvel-slurm-testing-slurm-00	*i-02028574274244586*	*t2.micro*
marvel-slurm-testing-gw	*i-01528621d1601b22d*	*t3a.small*
marvel-slurm-testing-slurmdb	*i-0361799692abc3dc4*	*t2.medium*

2. Calculate energy utilization for provisioned resources and show recommendations to deployment teams

(Legend: R denotes Recommendation Rank)

Under-Provisioned

Instance ID	Instance Type	R1	R2	R3
i-0011eee5f9bfa37d5	*t3.large*	*m7i-flex.large*	*m6i.large*	*m7i.large*
i-01e2bba24c61139a5	*t3.large*	*m7i-flex.large*	*m6i.large*	*m7i.large*
i-0378700bcf2ef68e7	*t3.large*	*m7i-flex.large*	*m6i.large*	*m7i.large*
i-061783c19995cf461	*t3.large*	*m7i-flex.large*	*m6i.large*	*m7i.large*
i-06b0d1885c8bfb73a	*t2.medium*	*c7i-flex.large*	*c6i.large*	*c5.large*
i-07d971daf49d26ba6	*t3.large*	*m7i-flex.large*	*m6i.large*	*m7i.large*
i-07fc1239a8432b1ff	*t2.nano*	*t3.small*	*t3.medium*	*t3.large*
i-0803db063f7902be8	*t3.small*	*t3.medium*	*m6i.large*	*t3.large*
i-0a30bd8f427568294	*t2.large*	*m7i.large*	*r7i.large*	*c7i-flex.xlarge*
i-0bcac5f25c482c094	*t3.large*	*m7i-flex.large*	*m6i.large*	*m7i.large*
i-0e7f58282fba77a57	*t3.xlarge*	*m7i-flex.xlarge*	*m6i.xlarge*	*m7i.xlarge*

Over-Provisioned

Instance ID	Instance Type	R1	R2	R3
i-07a7a6794367ded49	t3.xlarge	r7i.large	t3.xlarge	-
i-0f6c16ceb8ea482c2	t3.2xlarge	r6i.xlarge	r7i.xlarge	t3.2xlarge
i-0f97772fef03977df	t3.xlarge	r6i.large	r7i.large	t3.xlarge
i-0fda6069f43c06df1	t3.xlarge	r7i.large	t3.xlarge	-

3. Identify Unused Resources

Unused EC2 Resources

Instance Name	Instance ID	Instance Type	Unused Days
marvel-GenAI-partha	i-0a1748bca3199a63f	t3.medium	274
marvel-LLM-test	i-0b43ba541b65fcd8f	t3.2xlarge	175
marvel-mddp-devsecops	i-0c73c29f8e5d01576	t3.xlarge	162
marvel-EEaaS-Kubernetes	i-0ced2e371b3425fe5	t3.large	142
marvel-TEaas-openstack-kolla	i-0842262e5bb3e2915	t3.xlarge	141
marvel-awx-ansible	i-0d5da5ba883efcf5b	t3.xlarge	127
marvel-Neureda-registry	i-0dc340b462af4ed74	t3.medium	125
marvel-windows-jumboxtest	i-07a7a6794367ded49	t3.medium	122
marvel-PlatformEng-Keycloak	i-0161cf03311390911	t3.medium	121
marvel-PE-mddp-k8s-1	i-0bfe201057f0db5f9	t3.2xlarge	121
marvel-Prometheus	i-00f4b18e0c7615802	t2.micro	120
marvel-slurm-testing-slurmctl	i-02028574274244586	t3a.small	120
marvel-slurm-testing-slurm-00	i-0c10cab3a673568a5	t3a.small	120
marvel-slurm-testing-gw	i-01528621d1601b22d	t3a.small	120
marvel-slurm-testing-slurmdb	i-04cf6faf1252ba6a7	t3a.small	120
mfdm_qual_dd_pg	i-03b6cc42f6ee8679a	t2.medium	116
tcsnxgnmfdmins5	i-054d01dc2082a9968	t3.large	114

mfdm_qual_dd_pg	i-03b6cc42f6ee8679a	t2.medium	116
tcsnxgnmfdmins5	i-054d01dc2082a9968	t3.large	114
tcsnxgnpemonobs01	i-0a30bd8f427568294	t2.large	114
tcsnxgnpemonobs02	i-0644a1225a61cfad0	t3.medium	114
tcsnxgnpemonobs03	i-0f4a538834c7f8379	t2.large	114
tcsnxgnmfdmins1	i-06f35320739d13110	t3.medium	113
tcsnxgnmfdmins3	i-0dd30488163388001	t3.medium	113
tcsnxgnmfdmins2	i-030fdebbe195c6d05	t3.medium	113

4. Turn off workloads and node pools outside of business hours

Automate Shutdown/Startup :

What It Is: Scripted automatic shutdown/startup of workloads outside business hours.

Sustainability Benefits :

- Reduced energy consumption
- Lower greenhouse gas emissions
- Decreased e-waste from prolonged hardware lifespan

Identify Non-Essential Workloads :

What It Is: Analyzing workloads to determine which can be safely turned off during non-business hours.

Sustainability Benefits :

- Targeted energy reduction
- Minimized unnecessary resource utilization
- Enhanced overall efficiency

Autoscaling Configuration (Node Pools) :

What It Is: Configuring node pools to autoscale down to 0 nodes during non-business hours.

Sustainability Benefits :

- Dynamic energy consumption adjustment
- Significant reduction in idle resource energy waste
- Enhanced environmental responsiveness

Scheduled Node Pool Management :

What It Is: Scheduling node pool shutdowns/startups using platform-specific features or tools.

Sustainability Benefits :

- Predictable energy savings
- Reduced operational carbon emissions
- Improved resource utilization efficiency

Node Pool Sizing Optimization :

What It Is: Regular review to optimize node pool sizing.

Sustainability Benefits :

- Continuous energy efficiency improvement
- Reduced e-waste through minimized node replacements
- Lower environmental impact

5. Show energy efficient resource provisioning options during provisioning

Please refer to this _link_ for a guide on how to implement this.

6. Implement pre-checks for common issues in CI/CD pipeline to avoid failures in different stages of pipeline

Please refer to this _link_ for a guide on how to implement this.

Details :

1. Development (4/24)

Areas Implemented :

- Energy metrics for application using Intel RAPL, Kepler, Schaphandre
- Energy consumption of ML job, Model training
- Optimize source code for energy and carbon emissions using static code analysis
- Compare memory and energy utilization of docker and multi-stage docker files

Areas to be Implemented :

- Simplicity and Efficiency Trade-offs while developing algorithms
- Efficient Software Algorithms and data structures
- Efficient integration and delivery pipelines (Dev, Test, CI env)
- On demand development environments
- Minimizing unnecessary code execution
- Utilizing lazy loading
- Optimize energy consumption for Python and Java applications

- Leverage LLMs to suggest energy efficient algorithms for Java / Python applications in IDE
- Provide configurable deployment of applications with selected features to minimize resource utilization
- Calculate energy utilization at function level
- Optimize energy utilization of applications using AI/ML
- Memory optimization for applications
- Rightsizing containers and VM's based on resource utilization
- Compare and suggest open-source tools which are energy efficient
- Train models with custom rules for optimization of code
- Suggest optimized version of built-in functions which are energy efficient in IDE / static analyzer
- Code splitting - Convert code to smaller files which can be loaded on-demand
- Tree shaking: This technique helps reduce the overall application size by removing unused code from the final build, thus optimizing the web product and minimizing resource consumption
- Split large code repositories
- Remove dead code and unused data from application

2. Network (1/2)

Areas Implemented :

- Compare energy efficiency for synchronous (REST vs gRPC) and Asynchronous (Kafka vs RabbitMQ) messaging between microservices

Areas to be Implemented :

- Optimize data transfer (Implement local data processing, aggregation and compression techniques to avoid transfer of frequent, large amounts of data)

3. Deployment (7/19)

Areas Implemented :

- Tagging of resources to track usage
- Calculate energy utilization for provisioned resources and show recommendations to deployment teams
- Identify and share report on unused resources
- Turn off workloads and node pools outside of business hours
- Show energy efficient resource provisioning options during provisioning
- Implement pre-checks for common issues in CI/CD pipeline to avoid failures indifferent stages of pipeline
- Workload allocation on energy efficient servers, VMs

Areas to be Implemented :

- *Deploy application in a region that is closest to users*
- *Deploy in regions that are powered by renewables or energy efficient resources*
- *Energy efficient workload scheduling*
- *CPU frequency optimization for non-critical workloads*
- *Alerts during resource provisioning (to avoid over provisioning) about energy consumption and energy efficient options*
- *Calculate and minimize energy utilization per transaction*
- *Implement automated decision between horizontal vs vertical scaling of applications*
- *Set up minimal artifacts - Reduce disk, memory and processing demands by considering the software for a given purpose*
- *Schedule resource intensive tasks to execute when renewable energy source is available*
- *Consolidate application workloads to maximize server utilization and move idle servers to low power mode*
- *Measure (Idle/Normal/Peak) energy consumption of resources and identify scope for optimizations (Get power consumption per CPU core information from processor type)*
- *Identify unused resources in K8s using KOR tool*

4. Design (0/31)

Areas Implemented :

No Areas Implemented

Areas to be Implemented :

- *Sustainability as NFR*
- *Minimal Architecture (i.e. Developing minimal software for value)*
- *Replace long-running service with a simple Function as a Service*
- *Asynchronous processing and event-driven architectures*
- *API Design: protocol (TCP, UDP, custom wire protocols)*
- *API Design: data format used for request/response of your APIs: Protobuf/gRPC is better compared to JSON, XML*
- *API Design: Avro schemas instead of JSON for data serialization*
- *Reusable APIs vs P2P Integrations*
- *Choice of programming language: Go, Rust over Python, Java; Java 17 is better with respect to memory optimization*
- *Public cloud, choice of provider, and region*
- *Platform as a Service (PaaS) and Serverless (e.g., Function as a Service (FaaS))*
- *Containers and Kubernetes - Adopting energy efficient architectures such as containers or serverless*

- *Scheduling and batch vs. real-time - Demand Shaping*

- *Carbon awareness into the Kubernetes Scheduler*

- *Cost and Sustainability: Reducing operational costs and sustainability are aligned*

- *Enforce Quotas and Rate Limiting*

- *Reduce the Network Footprint*

- *Identify energy utilization of applications across different design choices*

- *Optimize data generation of applications*

- *Create reusable modules to minimize development and maintenance efforts*

- *Enable parallel processing using distributed computing and minimize energy consumption*

- *Measure KPIs on the features used by customers and remove unused features*

- *Energy saving mode for applications, system processes during periods of low load*

- *Energy budgets for applications*

- *Create Energy profile for application (Core Vs Non-core tasks - Logging, Monitoring, Audit, Redundancy etc...)*

- *Improve resilience of equipment to minimize redundancy (Ensure legacy hardware can support at least core features of software)*

- *Shared libraries for applications*

- *Server-side rendering (SSR): With SSR, a web page's HTML is generated on the server and sent to the client, resulting in quicker initial load times and improved search engine optimization*

- *React frameworks: Next.js as a framework has many built-in features for energy optimization*

- *Progressive Web Applications (Hybrid of native and web apps) are more energy efficient*

- *Optimize availability and scalability needs for application*

5. Storage (0/5)

Areas Implemented :

No Areas Implemented

Areas to be Implemented :

- *Dynamic provisioning of volumes for EC2 instance / Pod to automatically scale the storage size*

- *Identify and remove unused data and resources in cloud*

- *Implementing data retention, archival and deletion policies*

- *Utilize caching to store frequently used data*

- *Optimize log and metrics collection*

6. Quality (0/8)

Areas Implemented :

No Areas Implemented

Areas to be Implemented :

- *Code and design review for sustainability*
- *Document decisions using Key Design Decisions (KDDs) or Architectural Decision Records (ADRs)*
- *Profiling tools and static analysis - Code profiling and optimization tools*
- *Upgrade Runtimes and Modules*
- *Change aware testing*
- *Test Case Optimization*
- *Risk Based Testing*
- *On-demand and scalable Test Environments*

7. Operations (0/12)

Areas Implemented :

No Areas Implemented

Areas to be Implemented :

- *Automated sustainability maturity assessment for applications*
- *Observability and CarbonOps*
- *Automation*
- *Switching instances to chipsets that offer the same processing power at lower levels of energy consumption*
- *Sustainable DevOps practices integrate sustainability into the software delivery process*
- *Optimize resource allocation using ML*
- *Maximize utilization of resources*
- *Spot Instances*
- *Setup sustainability goals for infrastructure, development and applications and monitor them using dashboards and reports*
- *Reduce the cooling requirements of servers by ensuring that applications (servers) run in optimal temperature range*
- *Identify zombie workloads that consume resources (e.g., Application not running within container / VM)*
- *Dashboard for resource utilization at cluster, node, and pod level and grouped by applications in Kubernetes*